

# GraphRAG with Neo4j and LLMs

Prerequisite Concepts and Technologies



# Section 1

## Knowledge Graphs


# What is a Knowledge Graph?

# Definition

A knowledge graph is a structured representation of information that captures entities and their relationships in a network format.

## Key Characteristics

- Entities represented as nodes
  - Relationships as edges
  - Properties on both nodes and edges
  - Semantic meaning encoded in structure



# Knowledge Graph Components

## Nodes (Entities)

Represent real-world objects

Examples: Person, Movie,  
Company

## Edges (Relationships)

Connect entities meaningfully

Examples: ACTED\_IN,  
WORKS\_FOR

## Properties (Attributes)

Store additional information

Examples: name, age, date

**Example:** (Actor)-[ACTED\_IN {role: "Neo"}]->(Movie)

# Graph vs. Relational Databases

## Relational (SQL)

- Tables with rows and columns
- Joins for relationships
- Schema-rigid structure
- Optimized for transactions
- Slower for deep relationships

## Graph (Neo4j)

- Nodes and relationships
- Direct connections (no joins)
- Flexible schema
- Optimized for connections
- Fast multi-hop queries

# Knowledge Graph Use Cases

## Search & Discovery

Google, Amazon product graphs

## Fraud Detection

Financial transaction networks

## Social Networks

Friend connections, content graphs

## Healthcare

Drug interactions, patient records

## AI & ML

Context for LLMs, recommendations

## Enterprise

Data lineage, organizational charts

# Section 2

Neo4j Database

# Introduction to Neo4j

Neo4j is the world's leading **graph database** platform with native graph storage and processing.

## Key Features

- ACID compliant transactions
- Cypher query language
- Scalable to billions of nodes
- Graph algorithms library
- Cloud-native (AuraDB)



Graph Native

# Graph Database Concepts

## Index-Free Adjacency

Each node directly references its neighbors, eliminating the need for index lookups during traversal

### Labels

Categorize nodes (e.g., :Person, :Movie)

### Relationship Types

Define connection types (e.g., ACTED\_IN)

**Each node maintains a constant-time pointer list to adjacent edges and connected nodes.**

**Therefore, traversal from one node to another is O(1) (constant time) per relationship.**

# Neo4j Architecture Layers

## Cypher Query Language

Declarative graph query language

## Graph Engine

Query execution and transactions

## Page Cache

In-memory caching

## Native Graph Storage

Persistent storage

# Neo4j Architecture Layers

Layer	Role	Analogy
<b>Cypher Query Language</b>	User-facing interface	How you speak to the graph
<b>Graph Engine</b>	Executes Cypher, manages transactions	The CPU
<b>Page Cache</b>	Keeps hot data in memory	The RAM
<b>Native Graph Storage</b>	Stores the full graph on disk	The Hard Drive

# Neo4j AuraDB Cloud

Fully managed cloud database service for Neo4j - perfect for labs and production



## Cloud Native

No installation required



## Secure

Encrypted connections



## Fast

SSD-backed storage

## Free Tier Available!

Perfect for learning and development

# Section 3

Cypher Query Language

# Cypher Query Language

Cypher is Neo4j's **declarative graph query language** - like SQL for graphs but with ASCII-art syntax

## Why Cypher?

- Intuitive visual syntax
- Pattern matching
- Expressive and readable
- Optimized for graphs

// Example Pattern

(actor)-[:ACTED\_IN]->(movie)

Nodes in parentheses

Relationships in brackets

# Basic Cypher Syntax

```
// CREATE: Add nodes and relationships  
CREATE (n:Person {name: 'Alice'})
```

```
// MATCH: Find patterns in the graph  
MATCH (p:Person) RETURN p.name
```

```
// WHERE: Filter results  
MATCH (p:Person) WHERE p.age > 21
```

```
// WITH: Chain operations  
MATCH (p) WITH p.name AS name RETURN name
```

# Common Cypher Patterns

Find all actors in a movie

```
MATCH (a:Actor)-[:ACTED_IN]->(m:Movie {title: 'Inception'})  
RETURN a.name
```

Find shortest path between nodes

```
MATCH p=shortestPath((a:Person)-[*]-(b:Person))  
RETURN p
```

Aggregate and count

```
MATCH (d:Director)-[:DIRECTED]->(m:Movie)  
RETURN d.name, count(m) AS movies
```

# Section 4

Large Language Models

# Large Language Models (LLMs)

LLMs are neural networks trained on vast text corpora to understand and generate human language

## Scale

Billions of parameters

## Training

Massive text datasets

## Versatility

Multi-task capable

**Examples:** GPT-4, Llama, Mistral, Claude, Gemini

# Transformer Architecture

The transformer architecture revolutionized NLP with its attention mechanism

## Key Components

- Self-attention layers
- Positional encodings
- Feed-forward networks
- Layer normalization

## Attention Is All You Need

Vaswani et al., 2017

### Benefits:

- Parallel processing
- Long-range dependencies
- Scalable architecture

# Hugging Face Ecosystem

The leading platform for machine learning models and datasets

## Model Hub

500,000+ pre-trained models

## Transformers Library

Easy model deployment

## Datasets

Curated dataset library

## Inference API

Hosted model endpoints

# Popular Open-Source LLMs

## Mistral 7B

Powerful 7B parameter model

## Llama 2

Meta's open model family

## Gemma

Google's lightweight models

## Phi-3

Microsoft's efficient models

# LLM Capabilities

## Text Generation

Creative writing, code

## Question Answering

Information retrieval

## Summarization

Condensing information

## Translation

Multilingual support

## Classification

Sentiment, categorization

## Reasoning

Logic and inference

# Section 5

## Vector Embeddings

# What are Embeddings?

Embeddings are **dense vector representations** of text that capture semantic meaning in continuous space

## Characteristics

- Fixed-length vectors (e.g., 384, 768, 1536 dimensions)
- Semantically similar texts have similar vectors
- Learned from large text corpora
- Enable mathematical operations on meaning

Text

"knowledge graph"



Vector

[0.23, -0.15, 0.89, ...]

# Vector Space Properties

## Proximity = Similarity

Vectors close together in space represent similar concepts

### Example Cluster

"king" ≈ "queen" ≈ "monarch"

"dog" ≈ "puppy" ≈ "canine"

### Vector Operations

king - man + woman ≈ queen

Arithmetic on meaning!

# Semantic Similarity

Measure similarity between vectors using distance metrics

## Cosine Similarity

Measures angle between vectors (-1 to 1)

Most common for text embeddings

## Euclidean Distance

Straight-line distance in vector space

**Higher similarity = More relevant results**

# Popular Embedding Models

## **all-MiniLM-L6-v2**

Fast and lightweight, great for prototypes

## **text-embedding-ada-002**

OpenAI's production embedding model

## **BGE-large-en**

State-of-the-art open-source model

# Section 6

Retrieval-Augmented Generation

# Retrieval-Augmented Generation (RAG)

RAG enhances LLM outputs by retrieving relevant information before generation

## 1. Retrieve

Find relevant documents or data


## 2. Augment

Add retrieved context to prompt

## 3. Generate

LLM creates informed response

# RAG System Architecture



# Why Use RAG?

## Reduces Hallucinations

Grounds responses in facts

## Up-to-date Information

Access latest data

## Domain-Specific

Custom knowledge bases

## Source Attribution

Traceable answers

RAG > Fine-tuning for many use cases

# Vector Search in RAG

## How It Works

- Convert query to embedding
- Compare with stored vectors
- Return top-k similar items
- Use as context for LLM

**Advantage:** Captures semantic meaning beyond keywords

## Vector Databases

- Neo4j (with vector index)
- Pinecone
- Weaviate
- Qdrant
- Milvus

# Classical RAG Limitations

Traditional RAG systems face several challenges with complex queries

## Isolated Information

Documents retrieved independently without understanding connections

## No Relationship Context

Cannot traverse entity relationships or multi-hop reasoning

## Limited Semantic Understanding

Struggles with complex queries requiring domain knowledge

# Full-Text Search

Traditional keyword-based search using lexical matching

## Best For

- Exact phrase matching
- Names and identifiers
- Dates and numbers
- Domain-specific terms

## Limitations

- No semantic understanding
- Misses synonyms
- Word order matters
- Less flexible

Complements vector search for comprehensive retrieval

# Section 7

GraphRAG

# What is GraphRAG?

**GraphRAG = Knowledge Graphs + Retrieval-Augmented Generation**

Combines the structured relationships of knowledge graphs with the semantic understanding of LLMs for more accurate and contextual retrieval

## Key Innovation

Leverages both graph structure and vector embeddings to retrieve richer, more connected information

# GraphRAG: Deep Dive

## How It Works

- User query converted to embedding
- Vector search finds relevant nodes
- Graph traversal explores relationships
- Full-text search captures exact terms
- Context assembled from graph structure
- LLM generates informed response

## Key Advantages

- Preserves entity relationships
- Enables multi-hop queries
- Combines multiple retrieval methods
- Provides explainable results
- Better contextual understanding

# GraphRAG: Deep Dive

## GraphRAG Pipeline

Combining **vector search**, **full-text search**, and **graph traversal** to retrieve rich context for LLM-based reasoning.



### Vector Search

Find semantically similar nodes using embeddings.



### Full-Text Search

Match exact names, IDs, or keywords for precision.



### Graph Traversal

Follow relationships to gather connected facts.



### LLM Generation

Generate grounded answers from structured context.

GraphRAG = Vector Search + Full-Text Search + Graph Traversal → Context → LLM → Explainable Answer

# GraphRAG vs. Traditional RAG

## Traditional RAG

- Document-based retrieval
- Flat vector search
- Limited relationship context
- Isolated information

## GraphRAG

- Entity & relationship aware
- Graph traversal + vectors
- Rich contextual connections
- Multi-hop reasoning

# Hybrid Retrieval Strategy

Combining multiple search methods for optimal results

## Vector Search

Semantic similarity

## Full-Text Search

Keyword matching


## Graph Traversal

Relationship exploration

# Hybrid Retrieval: The Complete Picture

## Hybrid Retrieval in GraphRAG

Example query: "What proteins interact with drugs that treat hypertension?"



Hybrid Retrieval = Vector Search + Full-Text Search + Graph Traversal → LLM → Explainable, Connected Answer

# Benefits of GraphRAG

## Better Accuracy

More relevant context through relationships

## Rich Context

Connected information, not isolated facts

## Complex Queries

Handle multi-hop reasoning

## Explainable

Traceable through graph structure

The future of intelligent QA systems

# Section 09

Putting It All Together

# GraphRAG Best Practices

## ✓ Design Good Schema

Clear labels, meaningful relationships

## ✓ Create Indexes

Vector and full-text for performance

## ✓ Use Quality Embeddings

Choose appropriate model size

## ✓ Optimize Cypher

Profile queries, use LIMIT wisely

## ✓ Test Retrievers

Evaluate precision and recall

## ✓ Monitor Performance

Track latency and accuracy

# Common Pitfalls to Avoid

## ⚠️ Missing Indexes

Slow queries without proper vector/fulltext indexes

## ⚠️ Over-retrieval

Fetching too much context overwhelms the LLM

## ⚠️ Poor Graph Design

Unclear relationships make traversal ineffective

## ⚠️ Ignoring Graph Structure

Using only vector search defeats the purpose of GraphRAG

# GraphRAG Real-World Applications

## Healthcare

Drug interactions, medical knowledge graphs, patient care paths

## Finance

Risk analysis, fraud detection, entity relationships

## Research

Scientific literature analysis, citation networks

## Enterprise

Knowledge management, organizational intelligence

## E-commerce

Product recommendations with relationship context

## Cybersecurity

Threat intelligence, attack pattern analysis

# GraphRAG Real-World Applications



## Empathize

### Who are your users?

- Students
- Doctors
- Researchers
- Travelers
- Entrepreneurs

### What do they need?

- Quick answers
- Connections between data
- Trusted sources

### Pain points

- Too much unstructured info
- Hard to find related facts
- Time-consuming search



## Define

### Problem Statement

- "Users need a way to \_\_\_ because \_\_\_"
- Example: "Researchers need a way to link articles and datasets because information is siloed."

### Knowledge Graph Domain

- Healthcare
- Education
- Tourism
- Law
- Mobility

### Entities & Relationships

- Drug–Disease
- Student–Course–Topic
- Place–Review–Activity



## Ideate

### Possible Solutions

- Chatbot interface
- Recommendation panel
- Interactive graph viewer

### Retrieval Strategies

- Vector search for semantic similarity
- Full-text for exact terms
- Graph traversal for multi-hop context

### LLM Outputs

- Summaries
- Explanations
- Recommendations
- Insights with citations

# Section 8

LangGraph Framework

# What is LangGraph?

LangGraph extends LangChain with stateful, cyclical workflows for building complex agent systems

## Cycles & Loops

Support for iterative and conditional workflows


## State Management

Persistent state across workflow steps

## When to Use LangGraph

- Multi-step reasoning with decision points
- Agent-based systems that need memory
- Complex workflows with conditional logic
- Systems requiring human-in-the-loop

# LangGraph: GraphRAG Workflow



*Conditional routing and loops enable complex reasoning*